Problem Set 2: Data Wrangling

Your Name Here

Background

Political advertising has traditionally been focused on the medium of television, but in recent cycles, online advertising has become much more popular. In this problem set, you will explore a dataset that has information on Facebook ad spending and impressions by candidates in the 2018 election cycle in the United States. The variables in this data are described below.

Name	Description
cand_id	unique identifier code for candidate
cand_name	full name of the candidate
cand_name_last	last name of the candidate
party	party affiliation of the candidate (R = Republican, D = Democrat)
office	office being sought by candidate
state	state in which the candidate is running
incumbency	incumbency status of candidate (incumbent, challenger, or open seat)
spend	estimated total spending on Facebook ads by candidate
impressions	estimated total impressions of Facebook ads
ad_tone_attack	proportion of FB ads that mention candidate's opponent only
ad_tone_promote	proportion of FB ads that mention candidate only
ad_tone_contrast	proportion of FB ads that mention candidate and candidate's opponent

Question 1 (8 points)

Load the data using the `read_csv` function and save it as `fb_ads` (using this will automatically make `fb_ads` a tibble). In the text, describe how many candidates there are in the dataset.

Use `dplyr` functions to create a table with the number of candidates in each type of incumbency status in the data set. Save this table output as `incumbency_table` (for the autograder). Use the function `knitr::kable()` on this table to have a nicely formatted table produced in the knitted output.

Rubric: 2pt for loading the data (autograder); 1pt for describing the number of candidates (PDF); 3pts for creating the table (autograder); 2pt for using `kable()` to nicely format the output (PDF)

Answer 1

```r
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##     filter, lag
```

```
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

```r
library(readr)
library(knitr)
fb_ads <-
  read_csv("data/fb_ads.csv")
```

```
## Rows: 7014 Columns: 12
```

```
## -- Column specification -------------------------------------------------------------
## Delimiter: ","
## chr (7): cand_id, cand_name, cand_name_last, party, office, state, incumbency
## dbl (5): spend, impressions, ad_tone_attack, ad_tone_promote, ad_tone_contrast
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```r
fb_ads
```

```
## # A tibble: 7,014 x 12
##    cand_id  cand_name cand_name_last party office state incumbency  spend impressions ad_tone_attack
##    <chr>    <chr>     <chr>          <chr> <chr>  <chr> <chr>       <dbl>       <dbl>         <dbl>
## 1 ftm_223~ RAUNER, ~ RAUNER         R     Gover~ IL    Incumbent   2.88e5    18066608        0.413
## 2 ftm_223~ PRITZKER~ PRITZKER       D     Gover~ IL    Challenger  3.84e6    97795310        0.0210
## 3 ftm_223~ RAOUL, K~ RAOUL          D     Other~ IL    Open Seat   1.67e5     7246472.       0.0583
```

```
##  4 ftm_223~ HAROLD, ~ HAROLD          R      Other~ IL    Open Seat  3.85e4    4275971         0.0309
##  5 ftm_223~ WHITE, J~ WHITE           D      Other~ IL    Incumbent  7.07e3    1768474.        0.0367
##  6 ftm_223~ HELLAND,~ HELLAND         R      Other~ IL    Challenger 2.54e3     262490.        0.169
##  7 ftm_223~ FRERICHS~ FRERICHS        D      Other~ IL    Incumbent  1.38e5    9123175         0.0848
##  8 ftm_223~ SALAZAR,~ SALAZAR         D      Other~ CO    Open Seat  4.95e1       3000.        0.0171
##  9 ftm_223~ WEISER, ~ WEISER          D      Other~ CO    Open Seat  7.53e4    4722822.        0.0151
## 10 ftm_223~ HILL, ST~ HILL            R      Other~ CO    Open Seat  1.39e3      46994.        0.0180
## # i 7,004 more rows
## # i 2 more variables: ad_tone_promote <dbl>, ad_tone_contrast <dbl>
```

```r
incumbency_table <- fb_ads|>
  group_by(incumbency) |>
  count() |>
  ungroup()
knitr::kable(incumbency_table)
```

incumbency	n
Challenger	2510
Incumbent	2022
Open Seat	2482

Question 2 (7 points)

Filter the data to just US House and US Senate races and use this to create a tibble called `party_incumbent_promote` that has 6 rows that summarizes the average of `ad_tone_promote` for each combination of `party` and `incumbency`. Call the variable summarizing the promote variable as `promote_prop` and be sure to remove any missing values when computing the averages.

Use `knitr::kable()` to produce a nicely formatted table. In this call, set the `digits` arguments to 3 and use the `col.names` argument to pass a nicer set of names. You can use the following as a template:

```r
knitr::kable(my_table, col.names = c("Variable 1", "Variable 2", ...))
```

In the writeup, describe which type of candidate sponsored the most promoting ads on average.

Rubric: 3pts for creating `party_incumbent_promote` correctly (autograder); 2pt for a nicely formatted table (PDF); 1pt for changing the column names of the output table (PDF); 1pt for correctly identifying the type of candidate with highest average (PDF)

Answer 2

```r
library(dplyr)
library(knitr)
library(readr)
filtration <- fb_ads |>
  filter(office %in% c("US House", "US Senate"))
filtration

party_incumbent_promote <-
  filtration |>
  group_by(party, incumbency) |>
  summarize(promote_prop = mean(ad_tone_promote, na.rm = TRUE)) |>
  ungroup()
party_incumbent_promote

knitr::kable(party_incumbent_promote
             , col.names = c("Party", "Incumbency Status", "Promote Proportion"),
             digits = 3)
```

Democratic Incumbents promoted the most ads, on average.

Question 3 (7 points)

Create a new variable called `impressions_millions` that is the total Facebook ad impressions in millions (as opposed to single impressions). Make sure to save the resulting dataset back as `fb_ads`.

Create a histogram of this variable for just the US House races. Save the ggplot output as `plot_q3` and also print it to produce a plot in the output. In the text, describe the shape of the histogram and tell the reader if most of the House candidates had more than 10 million ads impressions on Facebook.

Rubric: 2pt for creating the new variable (autograder); 3pts for creating the histogram object (autograder); 2pts for answering the question about the histogram (PDF)

Answer 3

```
library(dplyr)
library(ggplot2)
fb_ads <- fb_ads |>
  mutate(impressions_millions = impressions / 100000)
fb_ads
```

```
## # A tibble: 7,014 x 13
##    cand_id  cand_name cand_name_last party office state incumbency  spend impressions ad_tone_attack
##    <chr>    <chr>     <chr>          <chr> <chr>  <chr> <chr>       <dbl>      <dbl>        <dbl>
##  1 ftm_223~ RAUNER, ~ RAUNER         R     Gover~ IL    Incumbent   2.88e5   18066608       0.413
##  2 ftm_223~ PRITZKER~ PRITZKER       D     Gover~ IL    Challenger  3.84e6   97795310       0.0210
##  3 ftm_223~ RAOUL, K~ RAOUL          D     Other~ IL    Open Seat   1.67e5    7246472.      0.0583
##  4 ftm_223~ HAROLD, ~ HAROLD         R     Other~ IL    Open Seat   3.85e4    4275971       0.0309
##  5 ftm_223~ WHITE, J~ WHITE          D     Other~ IL    Incumbent   7.07e3    1768474.      0.0367
##  6 ftm_223~ HELLAND,~ HELLAND        R     Other~ IL    Challenger  2.54e3     262490.      0.169
##  7 ftm_223~ FRERICHS~ FRERICHS       D     Other~ IL    Incumbent   1.38e5    9123175       0.0848
##  8 ftm_223~ SALAZAR,~ SALAZAR        D     Other~ CO    Open Seat   4.95e1       3000.      0.0171
##  9 ftm_223~ WEISER, ~ WEISER         D     Other~ CO    Open Seat   7.53e4    4722822.      0.0151
## 10 ftm_223~ HILL, ST~ HILL          R     Other~ CO    Open Seat   1.39e3      46994.      0.0180
## # i 7,004 more rows
## # i 3 more variables: ad_tone_promote <dbl>, ad_tone_contrast <dbl>, impressions_millions <dbl>
```

```
house_races <- fb_ads |>
  filter(office == "US House")
house_races
```

```
## # A tibble: 1,054 x 13
##    cand_id  cand_name cand_name_last party office state incumbency  spend impressions ad_tone_attack
##    <chr>    <chr>     <chr>          <chr> <chr>  <chr> <chr>       <dbl>      <dbl>        <dbl>
##  1 H0AL020~ ROBY, MA~ ROBY           R     US Ho~ AL    Incumbent   2.57e4    1155494.      0.0639
##  2 H0AL051~ BROOKS, ~ BROOKS         R     US Ho~ AL    Incumbent   1.57e4    1223495       0.0112
##  3 H0AL070~ SEWELL, ~ SEWELL         D     US Ho~ AL    Incumbent   6.55e3     412495       0.00295
##  4 H0AR030~ WOMACK, ~ WOMACK         R     US Ho~ AR    Incumbent   9.35e3     807998.      0.00466
##  5 H0AZ012~ GOSAR, P~ GOSAR          R     US Ho~ AZ    Incumbent   3.67e4    2079995       0.00948
##  6 H0CA030~ BERA, AM~ BERA           D     US Ho~ CA    Incumbent   6.66e4    3180908       0.0188
##  7 H0CA100~ DESAULNI~ DESAULNIER     D     US Ho~ CA    Incumbent   2.25e3      49497       0.0192
##  8 H0CA191~ DENHAM, ~ DENHAM         R     US Ho~ CA    Incumbent   1.32e5    2822929       0.417
##  9 H0CA270~ SCHIFF, ~ SCHIFF         D     US Ho~ CA    Incumbent   2.10e5    7565082.      0.0228
```

5

```
## 10 H0CA321~ CHU, JUDY CHU          D     US Ho~ CA    Incumbent  4.95e1      3000.        0.000239
## # i 1,044 more rows
## # i 3 more variables: ad_tone_promote <dbl>, ad_tone_contrast <dbl>, impressions_millions <dbl>
```

```
plot_q3 <- ggplot(house_races, aes(x = impressions_millions)) +
  geom_histogram(binwidth = 1, fill = "blue", color = "black") +
  labs(title = "Distribution of Facebook Ad Impressions for US House Races",
       x = "Ad Impressions (in Millions)",
       y = "Count")
plot_q3
```


Distribution of Facebook Ad Impressions for US House Races
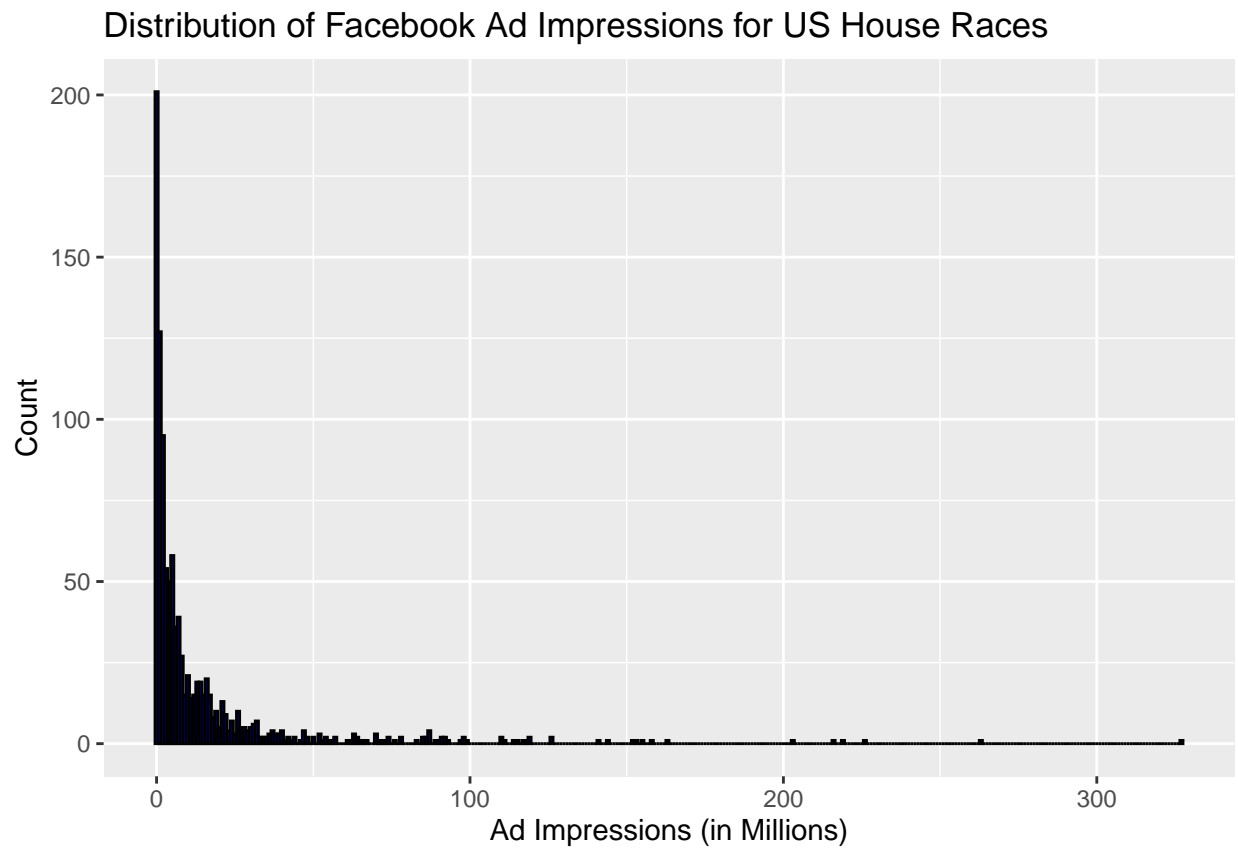
```
print(plot_q3)
```

Distribution of Facebook Ad Impressions for US House Races
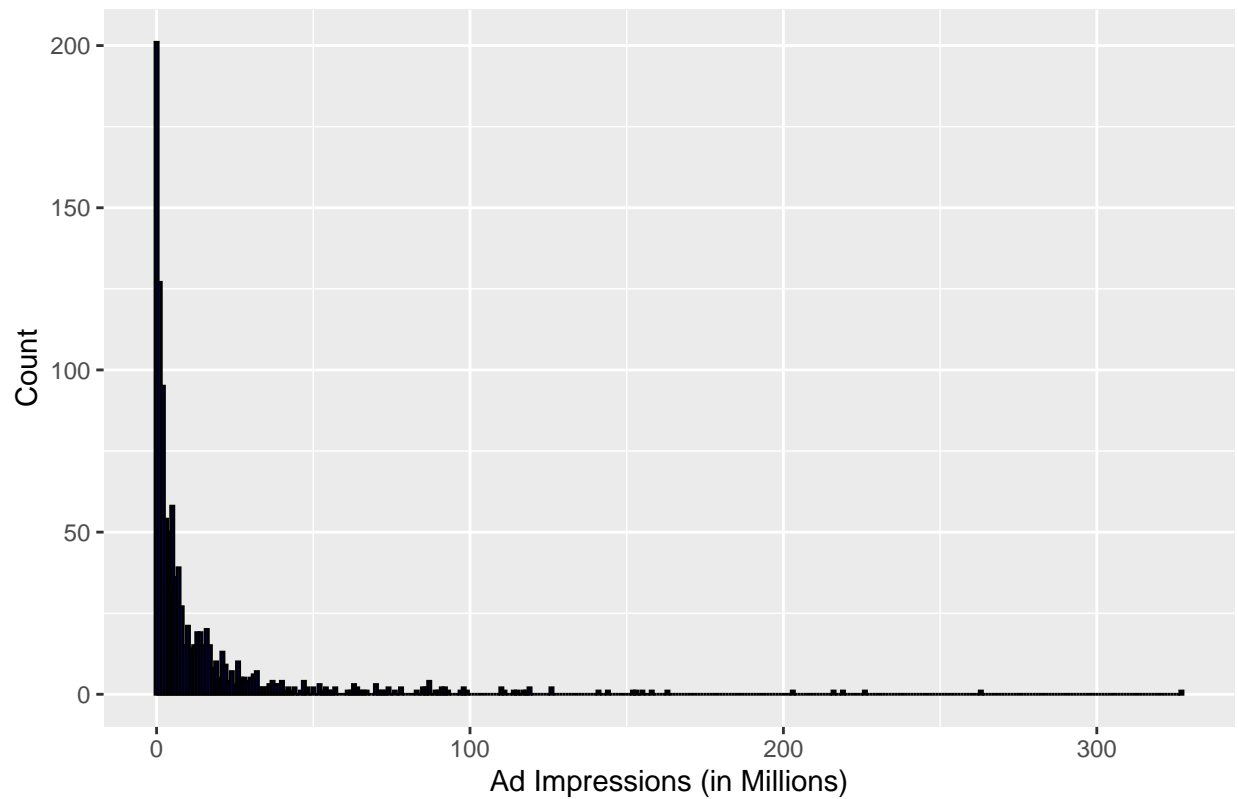
Most of the members of the House had under 10 million ad impressions. The result of this data being made into a histogram is therefore one where most House members fall close to the left-most end of the y-axis.

Question 4 (13 points)

Let's now recreate the following plot that shows the top 15 House candidates in terms of Facebook ad impressions.
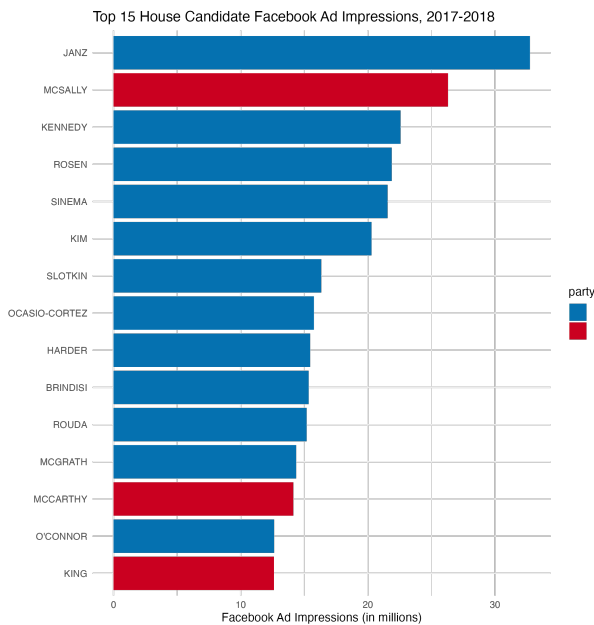
You should save the `ggplot` output as `fb_top_plot`. You should also write `fb_top_plot` on its own line in the chunk to produce the actual plot. The key features of this graph that you should replicate for the autograder are:

- The barplot should have candidate last names on the y-axis and the `impressions_millions` variable from question 3 on the x-axis.
- The data feeding into the `ggplot` call should only have US House candidates and only the candidates with the highest 15 `impressions_millions` values.
- The y-axis should be ordered in ascending values of `impressions_millions` so that the lowest values are at the bottom. You may want to manipulate `cand_name_last` to achieve this.
- The fill color of the bar plot should be mapped to the `party` variable (but not globally!).

You do not need to exactly match the labels, but you should have informative labels. The color does not need to match, but if you want to change the fill colors, you can use the `scale_fill_manual(values = c(R = "red", D = "blue"))` function (where you can change the red and blue to whatever you want).

Rubric: 3pts for correct axes (autograder); 3pts for correct data fed into `ggplot` (autograder); 3pts for the correct ordering of the y-axis (PDF); 3pts for fill being mapped to `party` (autograder); 1pt for plot being in knitted output and having informative labels (PDF)

Answer 4

```
library(dplyr)
library(ggplot2)
top_15_house <- fb_ads |>
  filter(office == "US House") |>
  arrange(desc(impressions_millions)) |>
```

```r
  slice_max(order_by = impressions_millions, n = 15)
top_15_house
```

```
## # A tibble: 15 x 13
##    cand_id  cand_name   cand_name_last party office state incumbency    spend impressions ad_tone_attack
##    <chr>    <chr>       <chr>          <chr> <chr>  <chr> <chr>         <dbl>       <dbl>         <dbl>
##  1 H8CA221~ JANZ, AN~   JANZ           D     US Ho~ CA    Challenger  1.24e6    32741906        0.0729
##  2 H2AZ081~ MCSALLY,~   MCSALLY        R     US Ho~ AZ    Incumbent   4.52e5    26300804        0.0459
##  3 H2MA040~ KENNEDY,~   KENNEDY        D     US Ho~ MA    Incumbent   3.20e5    22569831        0.0122
##  4 H6NV031~ ROSEN, J~   ROSEN          D     US Ho~ NV    Incumbent   6.88e5   21876048.        0.0651
##  5 H2AZ090~ SINEMA, ~   SINEMA         D     US Ho~ AZ    Incumbent   8.07e5   21550630.        0.0204
##  6 H8NJ032~ KIM, ANDY   KIM            D     US Ho~ NJ    Challenger  5.14e5   20286520.        0.182
##  7 H8MI081~ SLOTKIN,~   SLOTKIN        D     US Ho~ MI    Challenger  2.73e5    16337279        0.136
##  8 H8NY151~ OCASIO-C~   OCASIO-CORTEZ  D     US Ho~ NY    Challenger  3.17e5   15757652.        0.0398
##  9 H8CA101~ HARDER, ~   HARDER         D     US Ho~ CA    Challenger  2.74e5   15463694.        0.0862
## 10 H8NY221~ BRINDISI~   BRINDISI       D     US Ho~ NY    Challenger  2.63e5    15345345        0.0350
## 11 H8CA480~ ROUDA, H~   ROUDA          D     US Ho~ CA    Challenger  2.94e5   15186202.        0.0246
## 12 H8KY061~ MCGRATH,~   MCGRATH        D     US Ho~ KY    Challenger  2.38e5   14363196.        0.0181
## 13 H6CA221~ MCCARTHY~   MCCARTHY       R     US Ho~ CA    Incumbent   3.33e5    14142056        0.0211
## 14 H8OH122~ O'CONNOR~   O'CONNOR       D     US Ho~ OH    Open Seat   4.24e5    12631748        0.0316
## 15 H2NY030~ KING, PE~   KING           R     US Ho~ NY    Incumbent   1.68e5    12615799        0.0215
## # i 3 more variables: ad_tone_promote <dbl>, ad_tone_contrast <dbl>, impressions_millions <dbl>
```

```r
top_15_house <-
  top_15_house |>
  mutate(cand_name_last = factor(cand_name_last, levels = top_15_house$cand_name_last[order(top_15_house

fb_top_plot <- ggplot(top_15_house, aes(x = impressions_millions, y = cand_name_last, fill = party)) +
  geom_bar(stat = "identity") +
  labs(title = "Top 15 US House Candidates by Facebook Ad Impressions",
       x = "Ad Impressions (Millions)",
       y = "Candidate Last Name") +
  scale_fill_manual(values = c(R = "red", D = "blue")) +
  theme_minimal()

fb_top_plot
```

Top 15 US House Candidates by Facebook Ad Impressions

Code

```r
options(width = 100)
library(dplyr)
library(readr)
library(knitr)
fb_ads <-
  read_csv("data/fb_ads.csv")
fb_ads
incumbency_table <- fb_ads|>
  group_by(incumbency) |>
  count() |>
  ungroup()
knitr::kable(incumbency_table)
knitr::kable(my_table, col.names = c("Variable 1", "Variable 2", ...))
library(dplyr)
library(knitr)
library(readr)
filtration <- fb_ads |>
  filter(office %in% c("US House", "US Senate"))
filtration

party_incumbent_promote <-
  filtration |>
  group_by(party, incumbency) |>
  summarize(promote_prop = mean(ad_tone_promote, na.rm = TRUE)) |>
  ungroup()
party_incumbent_promote

knitr::kable(party_incumbent_promote
             , col.names = c("Party", "Incumbency Status", "Promote Proportion"),
             digits = 3)
library(dplyr)
library(ggplot2)
fb_ads <- fb_ads |>
  mutate(impressions_millions = impressions / 100000)
fb_ads

house_races <- fb_ads |>
  filter(office == "US House")
house_races

plot_q3 <- ggplot(house_races, aes(x = impressions_millions)) +
  geom_histogram(binwidth = 1, fill = "blue", color = "black") +
  labs(title = "Distribution of Facebook Ad Impressions for US House Races",
       x = "Ad Impressions (in Millions)",
       y = "Count")
plot_q3

print(plot_q3)

library(dplyr)
library(ggplot2)
```

```r
top_15_house <- fb_ads |>
  filter(office == "US House") |>
  arrange(desc(impressions_millions)) |>
  slice_max(order_by = impressions_millions, n = 15)
top_15_house

top_15_house <-
  top_15_house |>
  mutate(cand_name_last = factor(cand_name_last, levels = top_15_house$cand_name_last[order(top_15_hous

fb_top_plot <- ggplot(top_15_house, aes(x = impressions_millions, y = cand_name_last, fill = party)) +
  geom_bar(stat = "identity") +
  labs(title = "Top 15 US House Candidates by Facebook Ad Impressions",
       x = "Ad Impressions (Millions)",
       y = "Candidate Last Name") +
  scale_fill_manual(values = c(R = "red", D = "blue")) +
  theme_minimal()

fb_top_plot
```